

LIGN 167: Problem Set 1

October 3, 2023

Instructions

Collaboration policy: You may collaborate with up to three other students on this problem set (that is, total group size of 4). You should submit **one response** per group, using group submissions on Gradescope. When you submit your work, you must indicate who was in the group, and what each of your individual contributions were.

Starter code: We are providing starter code for the problem set. This starter code contains function signatures for all of the functions that you are defining. You **must** use this starter code when writing answers to the problems. If you do not follow the format of the starter code, **you will not receive credit**.

IMPORTANT: For each problem, you are expected to do 5 things:

1. Provide the correct answer to the problem.
2. Provide the solution that GPT-3.5 gives.
3. Provide the solution that GPT-4 gives.
4. If there are errors in the responses from either GPT-3.5 or GPT-4, explain where they occur.
5. For each problem, include links to your conversations with GPT-3.5 and GPT-4. See here for an example: <https://chat.openai.com/share/90efb80c-705e-4ed1-9e1e-68f13f4306cb>

Please submit this work in two files. One file should be named **ps1.py**. This should contain the correct answer to the problem. The second file should be named **gpt.py**. This should contain the code generated by GPT-3.5 and GPT-4, as well as your explanations of the mistakes that occurred.

In class we talked about *least squares regression*. In least-squares regression, we have a dataset that consists of two variables, $X = (x_1, \dots, x_n)$ and $Y = (y_1, \dots, y_n)$. We are trying to predict the values in Y from the values in X using the linear equation $y = a \cdot x + b$. More precisely, for each x_i , we use this equation to compute a predicted value:

$$\hat{y}_i = a \cdot x_i + b \tag{1}$$

Our goal is to minimize the total error of our predictions on the dataset. We want to find the values of a and b that minimize the quantity:

$$L = \sum_{i=1}^n (y_i - \hat{y}_i)^2 \quad (2)$$

$$= \sum_{i=1}^n (y_i - a \cdot x_i - b)^2 \quad (3)$$

In class, we found equations for the optimal values of the slope a and intercept b . These equations define *estimators* of the slope and intercept. The equations used the definition of the mean of a variable: $\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i$. The equations for the estimators of a and b that we found were:

$$a(x, y) = \frac{\sum_i x_i \cdot y_i - n\bar{x}\bar{y}}{\sum_i x_i^2 - n\bar{x}^2} \quad (4)$$

$$b(x, y) = \bar{y} - a(x, y) \cdot \bar{x} \quad (5)$$

Note that we are treating these estimators as functions: $a(x, y)$ takes in the values of x and y , and returns an estimated value of the slope, and similarly for $b(x, y)$.

Problem 1. Write a Python function `compute_slope_estimator`, which takes in two input variables, x and y . The variables x and y should be 1-dimensional NumPy arrays that have the same length n . The function should return the optimal value of the slope from Equation 4.

(If you are new to Python or Numpy, see this great tutorial: <http://cs231n.github.io/python-numpy-tutorial>. Also, please don't hesitate to come to office hours.)

Problem 2. Write a Python function `compute_intercept_estimator`, which takes in two input variables, x and y . The variables x and y should be 1-dimensional NumPy arrays that have the same length n . The function should return the optimal value of the intercept from Equation 5.

Problem 3. Write a function `train_model`, which takes in two 1-dimensional NumPy arrays of the same length, x and y . It should use `compute_slope_estimator` and `compute_intercept_estimator`, and return a tuple of values: the optimal value of the slope and the optimal value of the intercept.

The elements in the array y can be considered the labels in our training set: we use them to estimate the optimal values of the slope and intercept.

Problem 4. Write a function `dL_da` which takes four arguments: x_vals , y_vals , a , and b . The variables x_vals and y_vals are 1-dimensional NumPy arrays of length n . Your function should return the partial derivative $\frac{1}{n} \cdot \frac{\partial L}{\partial a}$:

$$\frac{1}{n} \cdot \frac{\partial L}{\partial a} = \frac{1}{n} \cdot \sum_{i=1}^n \frac{\partial f(x_i, y_i, a, b)}{\partial a} \quad (6)$$

In this equation, the function $f(x_i, y_i, a, b)$ is defined as:

$$f(x_i, y_i, a, b) = (y_i - a \cdot x_i - b)^2 \quad (7)$$

Problem 5. Write a function `dL_db` which takes four arguments: `x_vals`, `y_vals`, `a`, and `b`. The variables `x_vals` and `y_vals` are 1-dimensional NumPy arrays of length n . Your function should return the partial derivative $\frac{1}{n} \cdot \frac{\partial L}{\partial b}$:

$$\frac{1}{n} \cdot \frac{\partial L}{\partial b} = \frac{1}{n} \cdot \sum_{i=1}^n \frac{\partial f(x_i, y_i, a, b)}{\partial b} \quad (8)$$

The function f is defined in the previous problem.

Problem 6. Write a function `gradient_descent_step` which takes five arguments: `x_vals`, `y_vals`, `a`, `b`, and `k`. The variable `k` is a positive real number. This function will perform a single step of *gradient descent*. Given the input values of `a` and `b`, the function will compute updated values `a_updated` and `b_updated`:

$$a_{\text{updated}} = a - \frac{k}{n} \cdot \frac{\partial L}{\partial a} \quad (9)$$

$$b_{\text{updated}} = b - \frac{k}{n} \cdot \frac{\partial L}{\partial b} \quad (10)$$

In order to calculate the partial derivatives $\frac{\partial L}{\partial a}$ and $\frac{\partial L}{\partial b}$, you should call the functions that you wrote in Problems 4 and 5.

The function should return a tuple of the values `a_updated` and `b_updated`.

(Note that in the function signature provided for you, the variable `k` is set to a default value. Please leave this value unchanged.)

Problem 7. In the previous problem, you wrote a function which performs a single step of gradient descent. Each time a gradient descent step is performed, the parameter values `a` and `b` will improve by a small amount. As we will soon discuss in class, the loss function generally decreases after a gradient descent step, meaning that the model provides a better fit to the data.

The full gradient descent algorithm is simple: iteratively perform many gradient descent steps, getting improved parameter values with each step. More precisely, we start with initial parameter values `a_0` and `b_0`. At step t , we compute the updated parameters `a_t` and `b_t` by performing a gradient descent step on `a_{t-1}` and `b_{t-1}`, the parameters from step $t-1$.

Write a function `gradient_descent`, which takes five arguments: `x_vals`, `y_vals`, `a_0`, `b_0`, and `k`. `a_0` and `b_0` are initial parameter values for the algorithm. `k` is a positive integer, which is the number of gradient descent steps that should be performed. The function should implement the gradient descent algorithm, performing `k` gradient descent steps. It should return the final parameter values `a_k` and `b_k` as a tuple.

(Note that in the function signature provided for you, the variables `a_1`, `b_1`, and `k` are set to default values which you should leave unchanged.)

Problem 8. In the following problems, we will be looking at Einstein Summation (`np.einsum`). As we will discuss in class, Einstein Summation will allow us to greatly simplify the coding that we do later in the course.

Write a function that takes two arguments, `A` and `B`. Both `A` and `B` are 2-dimensional NumPy arrays of the same size. (A 2-dimensional array is a matrix. Each array has n rows

and m columns.) The function should return a 2-D NumPy array C of the same size, defined as follows:

$$C_{i,j} = A_{i,j} \cdot B_{i,j} \quad (11)$$

$C_{i,j}$ is the element of C in the i 'th row and j 'th column.

For example, suppose that the two input arrays are:

$$A = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}, \quad B = \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \end{bmatrix} \quad (12)$$

Then the output array will be:

$$C = \begin{bmatrix} 1 & 0 & 3 \\ 0 & 5 & 0 \end{bmatrix} \quad (13)$$

Your function must use `np.einsum` to perform this operation. See documentation here: <https://numpy.org/doc/stable/reference/generated/numpy.einsum.html>

Problem 9. Write a function that takes two arguments, A and B . A is a 2-D NumPy array, and B is a 1-D array. The size of A is (n,m) , while the size of B is (m) . (This means that matrix A has n rows and m columns, and vector B has m rows.) The function should return a 2-D array C of size (n,m) , defined as follows:

$$C_{i,j} = A_{i,j} \cdot B_j \quad (14)$$

For example, suppose that the two input arrays are:

$$A = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}, \quad B = \begin{bmatrix} 0 \\ 1 \\ 2 \end{bmatrix} \quad (15)$$

Then the output array will be:

$$C = \begin{bmatrix} 0 & 2 & 6 \\ 0 & 5 & 12 \end{bmatrix} \quad (16)$$

Your function must use `np.einsum` to perform this operation.

Problem 10. This example will involve 3-D NumPy arrays. The 3-D arrays we will be considering have size (b, n, m) . You can think of this as a collection of b matrices which each have size (n, m) . For example, here is a 3-D array with size $(2,3,4)$:

$$A = \begin{bmatrix} \begin{bmatrix} 1 & 2 & 3 & 4 \\ 5 & 6 & 7 & 8 \\ 9 & 10 & 11 & 12 \end{bmatrix} \\ \begin{bmatrix} 13 & 14 & 15 & 16 \\ 17 & 18 & 19 & 20 \\ 21 & 22 & 23 & 24 \end{bmatrix} \end{bmatrix}$$

As you can see, this 3-D array is made up of 2 matrices, each of size $(3,4)$.

Since there are 3 dimensions in the array, we need 3 indices in order to pick out a specific number. For example, $A_{0,2,3} = 12$, and $A_{1,0,2} = 15$. Make sure that you understand how this indexing works before moving on.

Write a function that takes two arguments, A and B. A is a 3-D NumPy array, and B is a 2-D array. The size of A is (b,n,m), and the size of B is (b,m). The function should return a 2-D array C of size (b,n), defined as follows:

$$C_{i,j} = \sum_k A_{i,j,k} \cdot B_{i,k} \quad (17)$$

This is batch matrix-vector multiplication, something that is used very often in deep learning.

Problem 11. Write a function that takes two arguments, A and B. A and B are both 3-D NumPy arrays. The size of A is (b,n,m), and the size of B is (b,m,p). The function should return a 3-D array C of size (b,n,p), defined as follows:

$$C_{i,j,q} = \sum_k A_{i,j,k} \cdot B_{i,k,q} \quad (18)$$

This is batch matrix-matrix multiplication, another common operation.